

Basic MLP with manually-derived Backprop

Teo Asinari

October 13, 2023

1 Introduction

Goal: To design, train and use a simple 3-layer MLP for binary classification of size-2 vectors.

Design: of the form

$$[(layer_size, Activation) \dots]$$


: [(2, ReLU), (2, ReLU), (1, Sigmoid)]

1.1 Diagrams

1.1.1 Vectorized Diagram (Equiv to Roger Grosse' 'Computational Graph')

$$\begin{array}{ccccccc} & w_1 & & w_2 & & w_3 & & t \\ & \searrow & & \searrow & & \searrow & & \searrow \\ x \rightarrow L_1 \rightarrow a_1 \rightarrow L_2 \rightarrow a_2 \rightarrow L_3 \rightarrow a_3 \rightarrow \mathcal{L} \end{array}$$

1.1.2 Expanded Diagram (Equiv. to Roger Grosse' 'Network Architecture')



1.2 Definitions

1.2.1 Remark on weight notation

$w_{i,j,k}$ is to say the weight at the i -th layer, j -th neuron, k -th weight. Hence w_{111} is the first weight of the first neuron in the first layer, etc.

1.2.2 Remark on layer notation

This is a sub-case of the weight notation. I.e., L_{ij} is the j -th neuron at the i -th layer, etc.

1.2.3 Neuron firing calculation

This is just a straightforward dot-product. We have:

$$L_{ij} = w_{ij}x_i$$

Where x_i in this case is referring to a more general notion of 'layer input', not necessarily just the first input to the network as in the diagrams above.

1.3 BackPropagation Derivation

Notation for derivative of loss w.r.t. to a function I will be using the following: $\bar{f} = \frac{\partial \mathcal{L}}{\partial f}$. This notation was introduced by Roger Grosse from the University of Toronto.

Pa(x) and Ch(x) these refer to the sets of parent and child vertices of a vertex in a graph.

General Approach Let's label the computational graph nodes as v_1, \dots, v_N with some topological ordering. Then, our general goal for backprop is to compute \bar{v}_i for $i \in 1, \dots, N$. With these, we can trivially calculate the weight updates. We compute a forward pass of the network, then set $v_N = 1$, then, for $i = N - 1, \dots, 1$, we have:

$$\bar{v}_i = \sum_{j \in Ch(v_i)} \bar{v}_j \frac{\partial v_j}{\partial v_i} \quad (\text{The Backprop Rule}) \quad (1)$$

1.3.1 Applying the backprop rule

Loss and final activation Then, going backwards through the 'computational graph', starting at the end:

$$\bar{\mathcal{L}} = 1 \quad (2)$$

$$\bar{a_3} = \bar{\mathcal{L}} \frac{\partial \mathcal{L}}{\partial a_3}$$

$$\begin{aligned}
\overline{a_3} &= (1) \frac{\partial \mathcal{L}}{\partial a_3} \\
\overline{a_3} &= \frac{\partial \mathcal{L}}{\partial a_3} \\
\overline{a_3} &= \frac{\partial}{\partial a_3} \frac{1}{2} (a_3 - t)^2 \\
\overline{a_3} &= (a_3 - t) \frac{\partial}{\partial a_3} (a_3 - t) \\
\overline{a_3} &= (a_3 - t) (1) \\
\overline{a_3} &= (a_3 - t)
\end{aligned} \tag{3}$$

Final layer *N.B.* I use σ to denote the sigmoid function here, not an activation function.

$$\begin{aligned}
\overline{L_3} &= \overline{a_3} \frac{\partial a_3}{\partial L_3} \\
\overline{L_3} &= \overline{a_3} \frac{\partial}{\partial L_3} \sigma(L_3) \\
\overline{L_3} &= \overline{a_3} \sigma(L_3) (1 - \sigma(L_3))
\end{aligned} \tag{4}$$

Final layer weights

$$\begin{aligned}
\overline{w_{3i}} &= \overline{L_3} \frac{\partial}{\partial w_{3i}} L_3 \\
\overline{w_{3i}} &= \overline{L_3} \frac{\partial}{\partial w_{3i}} \sum_j w_{3j} a_{2j} \\
\overline{w_{3i}} &= \overline{L_3} a_{2i}
\end{aligned} \tag{5}$$

Second layer activation

$$\begin{aligned}
\overline{a_{2i}} &= \overline{L_3} \frac{\partial}{\partial a_{2i}} L_3 \\
\overline{a_{2i}} &= \overline{L_3} \frac{\partial}{\partial a_{2i}} \sum_j w_{3j} a_{2j} \\
\overline{a_{2i}} &= \overline{L_3} w_{3i}
\end{aligned} \tag{6}$$

Second layer

$$\begin{aligned}\overline{L_2} &= \overline{a_2} \frac{\partial}{\partial L_2} a_2 \\ \overline{L_2} &= \overline{a_2} \frac{\partial}{\partial L_2} \text{ReLU}(L_2)\end{aligned}$$

Note that $d/dx(\text{ReLU}(x))$ is the heaviside step function $\theta(x)$:

$$\begin{aligned}\begin{cases} 1 & \text{if } x > 0 \\ 0 & \text{otherwise} \end{cases} \\ \overline{L_2} &= \overline{a_2} \theta(L_2) \cdot (1) \\ \overline{L_2} &= \overline{a_2} \theta(L_2)\end{aligned}\tag{7}$$

Second layer weights

$$\begin{aligned}\overline{w_{2i}} &= \overline{L_2} \frac{\partial}{\partial w_{2i}} L_2 \\ \overline{w_{2i}} &= \overline{L_2} \frac{\partial}{\partial w_{2i}} \sum_j w_{2j} a_{1j} \\ \overline{w_{2i}} &= \overline{L_2} a_{1i}\end{aligned}\tag{8}$$

Input layer activation

$$\begin{aligned}\overline{a_{1i}} &= \overline{L_2} \frac{\partial}{\partial a_{1i}} L_2 \\ \overline{a_{1i}} &= \overline{L_2} \frac{\partial}{\partial a_{1i}} \sum_j w_{2j} a_{1j} \\ \overline{a_{1i}} &= \overline{L_2} w_{2i}\end{aligned}\tag{9}$$

Input layer

$$\begin{aligned}\overline{L_1} &= \overline{a_1} \frac{\partial}{\partial L_1} a_1 \\ \overline{L_1} &= \overline{a_1} \frac{\partial}{\partial L_1} \text{ReLU}(L_1) \\ \overline{L_1} &= \overline{a_1} \theta(L_1) \cdot (1) \\ \overline{L_1} &= \overline{a_1} \theta(L_1)\end{aligned}\tag{10}$$

Input layer weights

$$\begin{aligned}\overline{w_{1i}} &= \overline{L_1} \frac{\partial}{\partial w_{1i}} L_1 \\ \overline{w_{1i}} &= \overline{L_1} \frac{\partial}{\partial w_{1i}} \sum_j w_{1j} x_{1j} \\ \overline{w_{1i}} &= \overline{L_1} x_{1i}\end{aligned}\tag{11}$$

1.4 Misc. Remarks

1.4.1 Rounding: Training vs Inference

Since we aim to train a binary classifier, the `round()` would be necessary for the correct output range. However since `round()` is not differentiable, we omit it during training, calculating fractional losses instead. We only include `round()` during inference.